

# birth-times series

January 23, 2017

```
In [1]: import pandas as pd
import numpy as np
import statsmodels.api as sm
from scipy import stats
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
In [2]: ls
```

```
README.md          images/          pair.ipynb
birth-times series.ipynb  individual.md   pair.md
data/              iterate.dat
```

```
In [3]: df = pd.read_csv('data/birth.txt')
df.head()
```

```
Out[3]:   num_births
0         295
1         286
2         300
3         278
4         272
```

```
In [4]: print df.shape[0]/31 #31 years
print df.shape
```

```
12
(372, 1)
```

```
In [5]: df['dates']=pd.date_range("1980-01-01", "2010-12-31", freq="1M")
```

```
In [6]: df.head()
```

```
Out[6]:   num_births      dates
0         295  1980-01-31
1         286  1980-02-29
2         300  1980-03-31
3         278  1980-04-30
4         272  1980-05-31
```

```
def acf_pacf(ts, lags): fig = plt.figure(figsize=(12,8)) ax1 = fig.add_subplot(211) fig =
sm.graphics.tsa.plot_acf(ts, lags=lags, ax=ax1) ax2 = fig.add_subplot(212) fig = sm.graphics.tsa.plot_pacf(ts,
lags=lags, ax=ax2)
```

```
In [7]: df['time'] = range(372)
```

```
In [8]: df['month'] = pd.DatetimeIndex(df.dates).month
df['year'] = pd.DatetimeIndex(df.dates).year
df['quarters'] = pd.DatetimeIndex(df['dates']).quarter
```

```
In [9]: df.head()
```

```
Out[9]:
```

|   | num_births | dates      | time | month | year | quarters |
|---|------------|------------|------|-------|------|----------|
| 0 | 295        | 1980-01-31 | 0    | 1     | 1980 | 1        |
| 1 | 286        | 1980-02-29 | 1    | 2     | 1980 | 1        |
| 2 | 300        | 1980-03-31 | 2    | 3     | 1980 | 1        |
| 3 | 278        | 1980-04-30 | 3    | 4     | 1980 | 2        |
| 4 | 272        | 1980-05-31 | 4    | 5     | 1980 | 2        |

```
In [10]: df = df.set_index('dates')
```

```
In [11]: df.head()
```

```
Out[11]:
```

|            | num_births | time | month | year | quarters |
|------------|------------|------|-------|------|----------|
| dates      |            |      |       |      |          |
| 1980-01-31 | 295        | 0    | 1     | 1980 | 1        |
| 1980-02-29 | 286        | 1    | 2     | 1980 | 1        |
| 1980-03-31 | 300        | 2    | 3     | 1980 | 1        |
| 1980-04-30 | 278        | 3    | 4     | 1980 | 2        |
| 1980-05-31 | 272        | 4    | 5     | 1980 | 2        |

```
In [12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 372 entries, 1980-01-31 to 2010-12-31
Data columns (total 5 columns):
num_births      372 non-null int64
time            372 non-null int64
month           372 non-null int32
year            372 non-null int32
quarters        372 non-null int32
dtypes: int32(3), int64(2)
memory usage: 13.1 KB
```

```
In [13]: sdf = df.groupby('month')['num_births'].mean()
print sdf.max()
print np.argmax(sdf)
```

```
334.161290323
8
```

```
In [14]: df_month = df.groupby('month').sum()
```

```
In [15]: df_month.head(12)
```


```
Out[15]:
```

|       | num_births | time | year  | quarters |
|-------|------------|------|-------|----------|
| month |            |      |       |          |
| 1     | 9493       | 5580 | 61845 | 31       |
| 2     | 8942       | 5611 | 61845 | 31       |
| 3     | 9577       | 5642 | 61845 | 31       |
| 4     | 9008       | 5673 | 61845 | 62       |
| 5     | 9315       | 5704 | 61845 | 62       |
| 6     | 9428       | 5735 | 61845 | 62       |

|    |       |      |       |     |
|----|-------|------|-------|-----|
| 7  | 10164 | 5766 | 61845 | 93  |
| 8  | 10359 | 5797 | 61845 | 93  |
| 9  | 10231 | 5828 | 61845 | 93  |
| 10 | 10008 | 5859 | 61845 | 124 |
| 11 | 9475  | 5890 | 61845 | 124 |
| 12 | 9719  | 5921 | 61845 | 124 |


```
In [16]: df_month['num_births'].plot(kind = 'line')
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x117550790>
```




```
In [17]: df_month['num_births'].plot(kind = 'area')
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x119ad1910>
```




```
In [18]: df_month['num_births'].plot(kind = 'bar')
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x119b01390>
```



```
In [19]: ax = plt.scatter(x = df_month.index, y =df_month.num_births)
```




```
In [20]: df_month
```

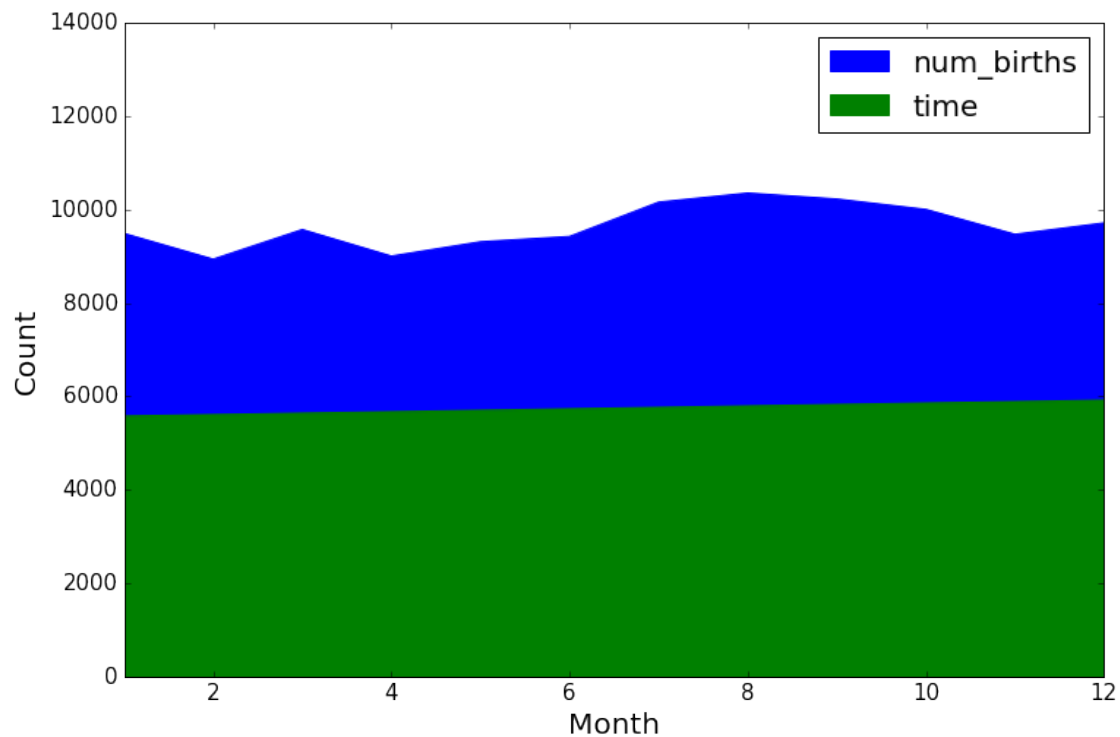
```
Out[20]:
```

|       | num_births | time | year  | quarters |
|-------|------------|------|-------|----------|
| month |            |      |       |          |
| 1     | 9493       | 5580 | 61845 | 31       |
| 2     | 8942       | 5611 | 61845 | 31       |
| 3     | 9577       | 5642 | 61845 | 31       |
| 4     | 9008       | 5673 | 61845 | 62       |
| 5     | 9315       | 5704 | 61845 | 62       |
| 6     | 9428       | 5735 | 61845 | 62       |
| 7     | 10164      | 5766 | 61845 | 93       |
| 8     | 10359      | 5797 | 61845 | 93       |
| 9     | 10231      | 5828 | 61845 | 93       |
| 10    | 10008      | 5859 | 61845 | 124      |
| 11    | 9475       | 5890 | 61845 | 124      |
| 12    | 9719       | 5921 | 61845 | 124      |

```
In [21]: ax = df_month['num_births'].plot('pie', figsize=(10,10),fontsize = 15)
ax.legend(fontsize = 10)
plt.show()
```



```
In [22]: ax = df_month['num_births'].plot('area', figsize=(20,10),fontsize = 15)
ax = df_month['time'].plot('area', figsize=(12,8),fontsize = 15)
ax.set_xlabel('Month', fontsize = 20)
ax.set_ylabel('Count',fontsize = 20)
ax.set_ylim(0,14000)
ax.legend(fontsize = 20)
plt.show()
```



```
In [23]: df_month
```

```
Out[23]:
```

|       | num_births | time | year  | quarters |
|-------|------------|------|-------|----------|
| month |            |      |       |          |
| 1     | 9493       | 5580 | 61845 | 31       |
| 2     | 8942       | 5611 | 61845 | 31       |
| 3     | 9577       | 5642 | 61845 | 31       |
| 4     | 9008       | 5673 | 61845 | 62       |
| 5     | 9315       | 5704 | 61845 | 62       |
| 6     | 9428       | 5735 | 61845 | 62       |
| 7     | 10164      | 5766 | 61845 | 93       |
| 8     | 10359      | 5797 | 61845 | 93       |
| 9     | 10231      | 5828 | 61845 | 93       |
| 10    | 10008      | 5859 | 61845 | 124      |
| 11    | 9475       | 5890 | 61845 | 124      |
| 12    | 9719       | 5921 | 61845 | 124      |

```
In [24]: idx = np.argmax(df_month.num_births)
print idx,df_month.num_births[idx]
df_month.loc[idx].num_births
```

```
8 10359
```

```
Out[24]: 10359
```

```
In [25]: df_month.idxmax()
```

```
Out[25]: num_births      8
time                  12
```

```
year          1
quarters      10
dtype: int64
```

```
In [26]: df_month.num_births.idxmax()
```

```
Out[26]: 8
```

```
In [27]: df_month.num_births[idx]
```

```
Out[27]: 10359
```

```
In [28]: df_year = df.groupby('year').sum()['num_births']
         indx = np.argmax(df_year)
         indx
```

```
Out[28]: 1993
```

```
In [29]: df.groupby('year').mean()['num_births'].idxmax()
```

```
Out[29]: 1993
```


```
In [30]: birthSeries = pd.Series(df['num_births'])
         bs = pd.Series(df['num_births'])
         birthSeries.head(2)
```

```
Out[30]: dates
         1980-01-31    295
         1980-02-29    286
         Name: num_births, dtype: int64
```

```
In [31]: ax = birthSeries.plot(figsize=(12,8), fontsize = 15,kind='line',grid=True)
         ax.set_xlabel('Dates',fontsize = 20)
         ax.set_ylabel('Birth Counts', fontsize = 20)
         ax.set_title('Birth Plot\n', fontsize = 30)
         plt.show()
```




## Birth Plot



Plot the data for 2006-2010, is the seasonal pattern more apparent?

```
In [32]: #birthSeries = pd.Series(df['num_births'])
         ax = birthSeries['2008':'2010'].plot(fontsize = 15, title='Birth Plot',figsize=(12,8))
         ax.set_xlabel('dates')
         ax.set_ylabel('Birth Counts')
         plt.show()
```



Use `df.resample('Q-NOV')` to get quarterly means that follow the seasons of the year (spring, summer, fall, winter).

```
In [33]: df['num_births'].head()
```

```
Out[33]: dates
1980-01-31    295
1980-02-29    286
1980-03-31    300
1980-04-30    278
1980-05-31    272
Name: num_births, dtype: int64
```

```
In [34]: bs.resample('A').mean().head() #year end
```

```
Out[34]: dates
1980-12-31    294.666667
1981-12-31    296.166667
1982-12-31    296.166667
1983-12-31    312.583333
1984-12-31    320.416667
Freq: A-DEC, Name: num_births, dtype: float64
```

```
In [35]: bs.resample('Q').mean().head()
```

```
Out[35]: dates
1980-03-31    293.666667
```

```


1980-06-30    272.666667
1980-09-30    314.000000
1980-12-31    298.333333
1981-03-31    289.000000
Freq: Q-DEC, Name: num_births, dtype: float64

```

```

In [36]: bs.plot(figsize=(12,8), label = 'monthly')
bs.resample('A').mean().plot(label = 'yearly')
bs.resample('Q-NOV').mean().plot(label = 'quarterly')
plt.legend();

```



```

In [37]: df['time^2'] = df['time']**2
df['time^3'] = df['time']**3
df['time^4'] = df['time']**4
df['time^5'] = df['time']**5
df.head()

```

```

Out[37]:
      num_births  time  month  year  quarters  time^2  time^3  time^4 \
dates
1980-01-31      295    0     1  1980         1        0        0        0
1980-02-29      286    1     2  1980         1        1        1        1
1980-03-31      300    2     3  1980         1        4        8       16
1980-04-30      278    3     4  1980         2        9       27       81
1980-05-31      272    4     5  1980         2       16       64      256

      time^5

```

```

dates
1980-01-31      0
1980-02-29      1
1980-03-31     32
1980-04-30    243
1980-05-31   1024

```

```
In [38]: type(df['num_births'])
```

```
Out[38]: pandas.core.series.Series
```

```
In [39]: type(df['num_births'].values)
```

```
Out[39]: numpy.ndarray
```

```
In [40]: y = df['num_births'].values
         X = df['time'].values
```

```
In [41]: model = sm.OLS(y, sm.add_constant(X)).fit()
```

```
In [42]: model.params
```


```
Out[42]: array([ 3.42255081e+02, -1.68099732e-01])
```

```
In [43]: model.fittedvalues[:10]
```


```
Out[43]: array([ 342.25508086,  342.08698113,  341.9188814 ,  341.75078167,
                341.58268193,  341.4145822 ,  341.24648247,  341.07838274,
                340.91028301,  340.74218328])
```

**0.0.1** If use `X = df['time']` instead of `df['time'].values`, then no need to add `index = df.index` since it's already indexed.

```
In [44]: bs.plot(figsize=(12,8))
         pd.Series(model.fittedvalues, index = df.index).plot();
         # df['time'].plot();
```



```
In [45]: y = df['num_births'].values #no index
         X = df['time'] #keep the index
         model = sm.OLS(y, sm.add_constant(X)).fit()
         bs.plot(figsize=(12,8))
         model.fittedvalues.plot(); # still indexed, same as bs
```



```
In [46]: model.fittedvalues[:10] #show the index
```

```
Out[46]: dates
1980-01-31    342.255081
1980-02-29    342.086981
1980-03-31    341.918881
1980-04-30    341.750782
1980-05-31    341.582682
1980-06-30    341.414582
1980-07-31    341.246482
1980-08-31    341.078383
1980-09-30    340.910283
1980-10-31    340.742183
dtype: float64
```

```
In [47]: model.summary()
```

```
Out[47]: <class 'statsmodels.iolib.summary.Summary'>
'''
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y    R-squared:                  0.264
Model:                            OLS    Adj. R-squared:            0.262
Method:                 Least Squares    F-statistic:                 133.0
Date:                   Mon, 23 Jan 2017    Prob (F-statistic):          1.72e-26
Time:                   22:10:59    Log-Likelihood:              -1794.5
No. Observations:          372    AIC:                        3593.

```

```

Df Residuals:      370    BIC:      3601.
Df Model:          1
Covariance Type:    nonrobust

```

|       | coef     | std err | t       | P> t  | [95.0% Conf. Int.] |
|-------|----------|---------|---------|-------|--------------------|
| const | 342.2551 | 3.125   | 109.535 | 0.000 | 336.111 348.399    |
| time  | -0.1681  | 0.015   | -11.531 | 0.000 | -0.197 -0.139      |

|                |       |                   |       |
|----------------|-------|-------------------|-------|
| Omnibus:       | 3.968 | Durbin-Watson:    | 0.316 |
| Prob(Omnibus): | 0.138 | Jarque-Bera (JB): | 3.093 |
| Skew:          | 0.099 | Prob(JB):         | 0.213 |
| Kurtosis:      | 2.600 | Cond. No.         | 428.  |

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""

```

```


In [48]: y = df['num_births'].values
        X = df[['time', 'time^2']]
        model = sm.OLS(y, sm.add_constant(X)).fit()
        bs.plot(figsize=(12,8))
        model.fittedvalues.plot()

```

```

Out[48]: <matplotlib.axes._subplots.AxesSubplot at 0x11b96fd10>

```



```
In [49]: model.summary()
```

```
Out[49]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.564
Model:                  OLS      Adj. R-squared:           0.562
Method:                 Least Squares    F-statistic:        239.0
Date:                   Mon, 23 Jan 2017    Prob (F-statistic):    2.62e-67
Time:                   22:10:59      Log-Likelihood:       -1697.0
No. Observations:       372      AIC:                  3400.
Df Residuals:           369      BIC:                  3412.
Df Model:                2
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [95.0% Conf. Int.]
-----
const          299.6021         3.600      83.232      0.000      292.524      306.680
time             0.5236         0.045     11.681      0.000         0.435         0.612
time^2          -0.0019         0.000    -15.941      0.000        -0.002        -0.002
=====
Omnibus:                 5.782    Durbin-Watson:           0.534
Prob(Omnibus):            0.056    Jarque-Bera (JB):         5.168
Skew:                     0.222    Prob(JB):                 0.0755
Kurtosis:                 2.630    Cond. No.                 1.84e+05
=====
```

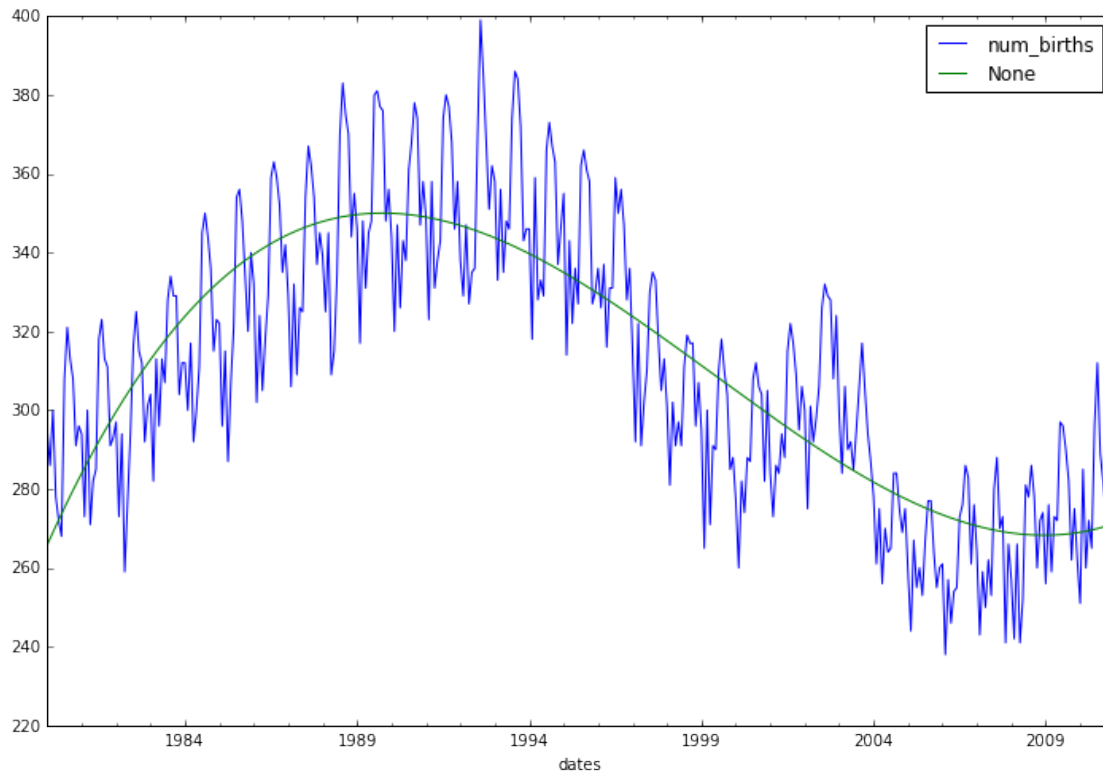
```
Warnings:
```

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.84e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
```

```
"""
```

```
In [50]: X = df[['time', 'time^2', 'time^3']]
model3 = sm.OLS(y, sm.add_constant(X)).fit()
bs.plot(figsize=(12,8))
model3.fittedvalues.plot()
plt.legend();
```





```
In [51]: model3.summary()
```

```
Out[51]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.701
Model:                        OLS      Adj. R-squared:           0.699
Method:                    Least Squares  F-statistic:                287.6
Date:                Mon, 23 Jan 2017    Prob (F-statistic):          4.31e-96
Time:                        22:11:00    Log-Likelihood:             -1627.0
No. Observations:                372     AIC:                       3262.
Df Residuals:                    368     BIC:                       3278.
Df Model:                        3
Covariance Type:                nonrobust
=====

```

|        | coef      | std err  | t       | P> t  | [95.0% Conf. Int.] |          |
|--------|-----------|----------|---------|-------|--------------------|----------|
| const  | 265.8165  | 3.963    | 67.074  | 0.000 | 258.023            | 273.609  |
| time   | 1.6238    | 0.093    | 17.529  | 0.000 | 1.442              | 1.806    |
| time^2 | -0.0093   | 0.001    | -15.996 | 0.000 | -0.010             | -0.008   |
| time^3 | 1.334e-05 | 1.03e-06 | 12.968  | 0.000 | 1.13e-05           | 1.54e-05 |

```

=====
Omnibus:                    3.849    Durbin-Watson:           0.778
Prob(Omnibus):              0.146    Jarque-Bera (JB):        3.070
Skew:                      0.106    Prob(JB):                0.215
=====


```

Kurtosis: 2.609 Cond. No. 7.67e+07  
=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 7.67e+07. This might indicate that there are strong multicollinearity or other numerical problems.  
"""

```
In [52]: X = df[['time', 'time^2', 'time^3', 'time^4']]
bs.plot(figsize = (12,8))
model4 = sm.OLS(y, sm.add_constant(X)).fit()
model4.fittedvalues.plot()
plt.legend();
```



```
In [53]: model4.summary()
```

```
Out[53]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

#### OLS Regression Results

```
=====
Dep. Variable:          y    R-squared:          0.726
Model:                OLS    Adj. R-squared:       0.723
Method:             Least Squares    F-statistic:        243.3
Date:                Mon, 23 Jan 2017    Prob (F-statistic):    8.01e-102
Time:                22:11:00    Log-Likelihood:       -1610.7
```

```

No. Observations:      372    AIC:      3231.
Df Residuals:          367    BIC:      3251.
Df Model:              4
Covariance Type:      nonrobust

```

|                | coef     | std err  | t                 | P> t  | [95.0% Conf. Int.] |           |
|----------------|----------|----------|-------------------|-------|--------------------|-----------|
| const          | 282.0876 | 4.719    | 59.782            | 0.000 | 272.809            | 291.367   |
| time           | 0.7359   | 0.177    | 4.164             | 0.000 | 0.388              | 1.083     |
| time^2         | 0.0015   | 0.002    | 0.781             | 0.435 | -0.002             | 0.005     |
| time^3         | -3.2e-05 | 7.87e-06 | -4.068            | 0.000 | -4.75e-05          | -1.65e-05 |
| time^4         | 6.11e-08 | 1.05e-08 | 5.810             | 0.000 | 4.04e-08           | 8.18e-08  |
| Omnibus:       |          | 3.293    | Durbin-Watson:    |       | 0.850              |           |
| Prob(Omnibus): |          | 0.193    | Jarque-Bera (JB): |       | 2.683              |           |
| Skew:          |          | 0.092    | Prob(JB):         |       | 0.261              |           |
| Kurtosis:      |          | 2.627    | Cond. No.         |       | 3.12e+10           |           |

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 3.12e+10. This might indicate that there are
strong multicollinearity or other numerical problems.
"""

```

```
In [54]: df.head()
```

```

Out[54]:
      num_births  time  month  year  quarters  time^2  time^3  time^4 \
dates
1980-01-31      295    0     1  1980         1         0         0         0
1980-02-29      286    1     2  1980         1         1         1         1
1980-03-31      300    2     3  1980         1         4         8        16
1980-04-30      278    3     4  1980         2         9        27        81
1980-05-31      272    4     5  1980         2        16        64       256

      time^5
dates
1980-01-31      0
1980-02-29      1
1980-03-31     32
1980-04-30    243
1980-05-31   1024

```

Now that you have fit trend, add in the monthly component via dummy variables to capture seasonality. You could also try to create a 'seasons of the year' variable and fit the quarterly time series instead of the original monthly time you plotted earlier...opportunity to play around.


```
In [55]: X = pd.get_dummies(df[['time', 'time^2', 'time^3', 'time^4', 'month']],
                             columns=['month']).drop('month_1', axis = 1)
```

```
In [56]: X.shape
```

```
Out[56]: (372, 15)
```

```
In [57]: model_m = sm.OLS(y, sm.add_constant(X)).fit()
         bs.plot(figsize=(12,8))
```

```
model_m.fittedvalues.plot(label = 'month_dummied')
plt.legend();
```



```
In [58]: model_m.summary()
```

```
Out[58]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y    R-squared:                  0.893
Model:                            OLS    Adj. R-squared:          0.888
Method:                           Least Squares    F-statistic:             197.2
Date:                            Mon, 23 Jan 2017    Prob (F-statistic):       4.32e-162
Time:                            22:11:01    Log-Likelihood:           -1436.6
No. Observations:                 372    AIC:                     2905.
Df Residuals:                     356    BIC:                     2968.
Df Model:                          15
Covariance Type:                  nonrobust
=====

```

|                   | coef       | std err  | t      | P> t  | [95.0% Conf. Int.] |           |
|-------------------|------------|----------|--------|-------|--------------------|-----------|
| const             | 278.7809   | 3.588    | 77.695 | 0.000 | 271.724            | 285.838   |
| time              | 0.6978     | 0.112    | 6.206  | 0.000 | 0.477              | 0.919     |
| time <sup>2</sup> | 0.0018     | 0.001    | 1.435  | 0.152 | -0.001             | 0.004     |
| time <sup>3</sup> | -3.254e-05 | 5e-06    | -6.505 | 0.000 | -4.24e-05          | -2.27e-05 |
| time <sup>4</sup> | 6.13e-08   | 6.69e-09 | 9.165  | 0.000 | 4.81e-08           | 7.45e-08  |

|          |          |       |        |       |         |         |
|----------|----------|-------|--------|-------|---------|---------|
| month_2  | -17.7778 | 2.988 | -5.950 | 0.000 | -23.654 | -11.902 |
| month_3  | 2.7017   | 2.988 | 0.904  | 0.366 | -3.174  | 8.578   |
| month_4  | -15.6583 | 2.988 | -5.240 | 0.000 | -21.535 | -9.782  |
| month_5  | -5.7612  | 2.988 | -1.928 | 0.055 | -11.638 | 0.115   |
| month_6  | -2.1232  | 2.988 | -0.711 | 0.478 | -8.000  | 3.754   |
| month_7  | 21.6105  | 2.988 | 7.231  | 0.000 | 15.733  | 27.488  |
| month_8  | 27.8915  | 2.989 | 9.332  | 0.000 | 22.014  | 33.769  |
| month_9  | 23.7519  | 2.989 | 7.946  | 0.000 | 17.874  | 29.630  |
| month_10 | 16.5464  | 2.989 | 5.535  | 0.000 | 10.668  | 22.425  |
| month_11 | -0.6604  | 2.990 | -0.221 | 0.825 | -6.540  | 5.219   |
| month_12 | 7.1958   | 2.990 | 2.407  | 0.017 | 1.315   | 13.076  |

|                |       |                   |          |
|----------------|-------|-------------------|----------|
| Omnibus:       | 8.418 | Durbin-Watson:    | 0.733    |
| Prob(Omnibus): | 0.015 | Jarque-Bera (JB): | 8.343    |
| Skew:          | 0.334 | Prob(JB):         | 0.0154   |
| Kurtosis:      | 3.303 | Cond. No.         | 7.93e+10 |


#### Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 7.93e+10. This might indicate that there are
strong multicollinearity or other numerical problems.
"""
```

Plot the dates variable (x) against the residuals (y) of the final model (including the seasonality term). Is there an obvious pattern of the residuals with respect to time? If there is any autocorrelation left in the model, there will be some pattern in your residual and we'll learn to address that in the afternoon.

```
In [59]: df['resid'] = model.resid
```

```
In [60]: df.plot(x = 'time', y = 'resid',figsize = (12,8), kind = 'scatter');
```



```
In [61]: pd.__version__
```

```
Out[61]: u'0.19.1'
```